

GazeSpeak: Academic Paper Questionnaire & Comprehensive Technical Answers

This document contains detailed, publication-ready answers for the **GazeSpeak** academic paper, derived from the codebase and technical design of the system.

1. Introduction and Problem Statement

1.1 Primary Problem and Target Audience

- bullet **Target Population:** Individuals suffering from severe motor disabilities, specifically **Amyotrophic Lateral Sclerosis (ALS)**, spinal cord injuries, brainstem stroke, or Locked-in Syndrome (LIS), who retain eye movement control but lose voluntary muscle movement and vocal ability.
- bullet **Primary Problem:** Existing Augmentative and Alternative Communication (AAC) devices rely either on high-cost specialized hardware (infrared gaze trackers like Tobii Dynavox costing 5,000–15,000) or complex multi-step eye-gesture coding schemes that carry a steep learning curve for patients and caregivers.
- bullet **GazeSpeak Solution:** A low-cost, software-only desktop AAC solution that enables high-accuracy eye-gaze typing and interactive conversation using a **standard commodity RGB webcam**, without specialized hardware or infrared illuminators.

1.2 Significance & Limitations of Existing Solutions

- bullet **High Financial Barrier:** Commercial eye-gaze AAC hardware requires specialized infrared illuminators and custom eye cameras, making them inaccessible in resource-constrained environments or developing regions.
- bullet **Physical Exhaustion & Eye Fatigue:** Traditional continuous point-and-dwell screen keyboards require precise 2D gaze tracking across high-density key matrices, causing rapid eye strain ("Midas Touch" problem where unintentional looks trigger unwanted keypresses).
- bullet **Setup Complexity:** Traditional 9-point or 16-point screen calibration requires sustained focal fixation, which is taxing for patients with nystagmus or fatigue.
- bullet **GazeSpeak Advances:**
 - 1. Horizontal 3-Point Calibration:** Simplifies setup to just 3 horizontal calibration steps (Left, Center, Right).
 - 2. Step-Based (Discrete) Gaze Navigation:** Replaces continuous 2D cursor placement with discrete step-based highlighting (looking Left/Right steps key highlights like arrow keys), drastically reducing required spatial gaze precision.
 - 3. Hysteresis & Neutral Dwell Arming:** Dwell timers only arm when gaze returns to the neutral Center zone, eliminating accidental activations while stepping.
 - 4. Hybrid LLM-Powered Prediction & Binary Conversation:** Combines offline fast dictionary prediction with Groq LLM (Llama 3.3 70B) for context-aware word prediction, abbreviation expansion, and binary (Left/Right choice) natural language dialogue.

1.3 Overarching Goal & Contributions

- bullet **Primary Objective:** Democratize accessible eye-gaze communication by converting ordinary consumer webcams into intuitive, low-latency AAC typing interfaces.
- bullet **Key Contributions:**

- 1. Step-Based Horizontal Gaze Ergonomics:** A novelty in low-cost AAC design that maps discrete horizontal eye movement (Left/Center/Right) to grid traversal and dwell selection.
 - 2. Hybrid Dual-Engine AAC Pipeline:** Real-time integration of offline frequency-ranked dictionary autocomplete with asynchronous cloud LLM (Groq Llama-3.3 70B) for phrase completion, context-aware abbreviation expansion (e.g., "ftk" → "from the kitchen"), and caretaker-guided binary decision trees.
 - 3. Multi-Modal Emergency SOS System:** Real-time Eye Aspect Ratio (EAR) rapid-blink monitoring (lge 5 blinks in 4 seconds) triggering a local audio siren (**winsound**) and automated caregiver SMS notification (Twilio API).
 - 4. Open-Source Implementation:** Built entirely in Python using open libraries (**PyQt6, OpenCV, MediaPipe Face Mesh, NumPy, pytttsx3, AssemblyAI**).
-

2. Background and Related Work

2.1 Technical Foundations & Research Areas

GazeSpeak synthesizes methodologies across three main research disciplines:

- 1. Computer Vision & Facial Landmark Detection:** Leverages Google MediaPipe Face Landmarker (478 3D landmarks, including 10 dedicated iris landmarks) to estimate iris position relative to eye corners in standard RGB video streams without IR lighting.
- 2. Augmentative and Alternative Communication (AAC):** Extends classic gaze-typing paradigms (e.g., Dasher, EyeSkills, GazeTalk) by prioritizing discrete directional movement over fine 2D spatial pointing.
- 3. Natural Language Processing & Adaptive Interfaces:** Integrates statistical language modeling (ngram/frequency dictionaries) and Large Language Models (LLMs) to maximize Keystroke Savings (KSS) and Words Per Minute (WPM).

2.2 Differentiation from Existing Systems

Feature / Dimension	Commercial AAC (e.g., Tobii Dynavox)	Conventional Webcam Gaze-Typing	GazeSpeak (Our System)
Hardware Required	IR illuminator + Dedicated Eye Camera (5,000–15,000)	Standard Webcam	Standard RGB Webcam (10–30)
Calibration Requirement	Complex 9 to 16-point 2D grid calibration	9-point focal calibration	Fast 3-point horizontal calibration (Left, Center, Right)
Navigation Paradigm	2D Continuous Gaze Cursor	2D Continuous Gaze Cursor	1D/2D Discrete Step-Based Highlight Traversal
Dwell Selection Safety	High rate of "Midas Touch" accidental triggers	High rate of accidental triggers	Safe Dwell Arming (Dwell arms ONLY upon returning to Center)
Conversation Paradigm	Manual Letter-by-Letter Typing	Manual Letter-by-Letter Typing	Interactive LLM Binary Choice (Left/Right) & Abbrev Expansion

Emergency SOS	Dedicated external hardware button	None or simple visual button	Automated Rapid-Blink EAR Detection + Siren + Twilio SMS
---------------	------------------------------------	------------------------------	--

2.3 Key Academic References & Influences

- 1. GazeSpeak (Microsoft Research / Purdue, 2017):** Inspired the fundamental concept of low-cost eye-gaze communication leveraging smartphone/webcam algorithms.
- 2. MediaPipe Face Mesh & Iris Tracking (Lugaresi et al., 2019 / Google Research):** Provides real-time 3D facial landmark detection and Sub-pixel iris center estimation (468-477 landmark indices).
- 3. Blink Detection via Eye Aspect Ratio (Soukupová & Mach, 2016):** Provided the mathematical foundation for real-time EAR calculation to detect blinks robustly across frames:
$$EAR = \frac{|p_2 - p_6| + |p_3 - p_5|}{2 |p_1 - p_4|}$$
- 4. Keystroke Savings & Predictive AAC (Trnka et al., 2009):** Guided the implementation of word prediction and abbreviation expansion to reduce physical patient fatigue.

3. System Design and Methodology

3.1 Architectural Breakdown

The system architecture follows a modular, signal-driven, multi-threaded layout implemented in Python and PyQt6:



Main System Modules:

- 1. Gaze Tracking Module (gaze/tracker.py):** Runs in a background **QThread**, capturing video frames, running MediaPipe Face Landmarker, computing normalized eye/iris ratios, performing exponential moving average smoothing, and evaluating EAR blinks.
- 2. Calibration Module (gaze/calibration.py):** Executes a 3-stage horizontal sequence (Look Left, Center, Right), filters outlier samples using Interquartile Range (IQR), and derives personalized horizontal gaze boundaries.

3. **UI & Keyboard Engine** (`ui/keyboard_widget.py`, `ui/prediction_bar.py`): Manages the QWERTY layout grid and suggestion strip. Handles discrete highlight navigation and visual dwell progress rings (`QPainter` rendering at ~60fps).
 4. **Word Prediction Engine** (`prediction/predictor.py`): A dual-tier predictor combining instant local dictionary lookups with asynchronous context-aware completions via the Groq API (`llama-3.3-70b-versatile`).
 5. **Abbreviation Expansion Engine** (`prediction/abbreviation_expander.py`): Evaluates short consonant-heavy tokens (e.g. "`ftk`") against caretaker context and conversation history to present full expanded options (e.g. "`from the kitchen`").
 6. **Binary Conversation Engine** (`prediction/conversation_engine.py`): Converts caretaker questions into binary decision trees, presenting two short contrasting options (`LEFT` card vs `RIGHT` card) that the patient chooses by looking left or right.
 7. **Emergency & SOS Alert Module** (`alerts/blink_alert.py`, `ui/emergency_screen.py`): Monitors consecutive rapid blinks, playing an immediate siren alarm (`winsound`) and dispatching SMS via Twilio.
 8. **Text-to-Speech & Speech-to-Text** (`ui/sentence_bar.py`, `audio/__init__.py`): Provides offline voice playback via `pyttsx3` and live caretaker voice transcription via AssemblyAI websockets.
-

3.2 Iris Tracking & Horizontal Eye Direction Algorithm

GazeSpeak utilizes MediaPipe's 478 face landmark mesh model:

bullet **Eye Corner Indices:**

bullet Left Eye: Inner corner = 133, Outer corner = 33, Top = 159, Bottom = 145

bullet Right Eye: Inner corner = 362, Outer corner = 263, Top = 386, Bottom = 374

bullet **Iris Landmark Indices:**

bullet Left Iris center = 468 (ring points 469-472)

bullet Right Iris center = 473 (ring points 474-477)

Gaze Ratio Calculation Procedure:

1. Extract pixel coordinates for iris center (p_{iris}), outer corner (p_{outer}), and inner corner (p_{inner}).

2. Calculate total eye width in pixels:

$$W_{\text{eye}} = |p_{\text{inner}} - p_{\text{outer}}|_2$$

3. Compute horizontal gaze ratio for each eye:

$$R_x = \frac{x_{\text{iris}} - x_{\text{outer}}}{W_{\text{eye}}}$$

4. Average ratios across left and right eyes to obtain a combined raw gaze metric:

$$\bar{R}_x = \frac{R_{x, \text{left}} + R_{x, \text{right}}}{2}, \quad \bar{R}_y = \frac{R_{y, \text{left}} + R_{y, \text{right}}}{2}$$

5. Confidence is estimated from bilateral eye agreement:

$$C = 1.0 - |R_{x, \text{left}} - R_{x, \text{right}}|$$

3.3 3-Point Horizontal Calibration Process

Calibration establishes the patient's individual eye range of motion without requiring complex 2D targets:

1. **Phases:** The user fixates on three visual target dots on the screen: **LEFT** ($x=8\%$ screen width), **CENTER** ($x=50\%$), and **RIGHT** ($x=92\%$).

2. Data Acquisition & Outlier Rejection:

- Captures 60 valid gaze frames (approx 2 seconds at 30 fps) per target point.

- Applies Interquartile Range (IQR) filtering to remove fixation drift and head jitter:

$$Q_1 = P_{\{25\}}(X), \quad Q_3 = P_{\{75\}}(X), \quad \text{IQR} = Q_3 - Q_1$$

$$\text{Keep } x_i \text{ if } x_i \in [Q_1 - 1.5 \cdot \text{IQR}, Q_3 + 1.5 \cdot \text{IQR}]$$

- Computes the mean gaze ratio ($\bar{R}_{\text{left}}$, $\bar{R}_{\text{center}}$, $\bar{R}_{\text{right}}$) from filtered samples.

3. Validation & Threshold Mapping:

- Validates that total span $|\bar{R}_{\text{right}} - \bar{R}_{\text{left}}| > 0.05$ and that $\bar{R}_{\text{center}}$ lies strictly between extreme bounds.

- Screen mapping scales raw ratio linearly:

$$S_x = \frac{\bar{R}_x - \bar{R}_{\text{left}}}{\bar{R}_{\text{right}} - \bar{R}_{\text{left}}}$$

- Calculates asymmetric zone boundaries around normalized center S_{center} :

$$\text{Zone}_{\text{LEFT}} < S_{\text{center}} - \max(0.55 \cdot S_{\text{center}}, 0.04)$$

$$\text{Zone}_{\text{RIGHT}} > S_{\text{center}} + \max(0.45 \cdot (1 - S_{\text{center}}), 0.12)$$

3.4 Step-Based Gaze Navigation & Hysteresis

Instead of free 2D cursor control, GazeSpeak categorizes horizontal eye direction into three discrete zones:

$$\text{Gaze Zone} = \begin{cases}$$

$$\text{LEFT} & \text{if } S_x < \text{Zone}_{\text{LEFT}} \\$$

$$\text{RIGHT} & \text{if } S_x > \text{Zone}_{\text{RIGHT}} \\$$

$$\text{CENTER} & \text{otherwise (neutral)}$$

$$\end{cases}$$

Navigation Dynamics:

- Hysteresis Filtering:** A candidate zone change must persist for $N = 2$ consecutive video frames (**HYSTERESIS_FRAMES = 2**) before the system commits to a zone transition. This prevents accidental stepping caused by momentary saccades or noise.

- Single-Step Traversal:** Transitioning from **CENTER** to **LEFT** steps the highlighted key leftward by one key; **CENTER** to **RIGHT** steps rightward by one key.

- Safety Dwell Arming:** Stepping disarms the selection dwell timer. Returning gaze to **CENTER** stops navigation and **arms** the dwell timer on the highlighted key.

3.5 Circular Navigation & Layout Wrapping

Key highlight navigation wraps seamlessly across keyboard rows and the word prediction bar:

- Moving right past the last key of a row wraps to column 0 of the next row down.

- Moving right past the bottom action row wraps up to the **Word Prediction Bar** (suggestion strip).

- Moving right past the last word prediction loops back to the top QWERTY row (**Row 0, Col 0**).

- Moving left past the top-left key wraps back up to the prediction bar.

3.6 Dwell Selection Mechanism

- bullet **Timer & Resolution:** Driven by a 60 fps Qt timer (`QTimer(interval=16ms)`).
- bullet **Dwell Accumulation:** When `_dwell_armed == True` and user gaze is steady at **CENTER**, progress accumulates:

$$\Delta P = \frac{16 \text{ ms}}{\text{DwellTime} - \text{target}}$$

- bullet **Visual Progress Ring:** A smooth arc is drawn around the target key using `QPainter.drawArc()`, spanning 0° to $360^\circ \times P$, accompanied by a radial gradient indicator at the leading edge.
 - bullet **Selection & Re-arming Cooldown:** Upon reaching $P = 1.0$, the key event fires, a visual green flash activates (`_selected_color`), and dwell is immediately **disarmed** for a 700 ms refractory period to prevent accidental double-selection.
-

4. Implementation Details and Tech Stack

4.1 GUI Framework: PyQt6

- bullet **Rationale:** Provides native C++ rendering speed, hardware-accelerated `QPainter` custom vector graphics, high-DPI scaling support, and a robust thread-safe Signal/Slot architecture.
- bullet **Key UI Components:**
 - bullet Custom custom-painted widgets (`KeyboardWidget`, `PredictionBar`, `ConversationPanel`, `QuickPhrasesPanel`) using anti-aliased `QPainter` vectors.
 - bullet Multi-page UI switching via `QStackedWidget` (Keyboard, Quick Phrases, Settings, Conversation Mode).
 - bullet Frameless full-screen overlays (`CalibrationScreen`, `EmergencyScreen`).

4.2 Computer Vision: OpenCV (cv2)

- bullet Captures 600x480 RGB frames from the system webcam at 30 fps (`cv2.VideoCapture`).
- bullet Performs horizontal mirror flipping (`cv2.flip(frame, 1)`) for intuitive natural eye control.
- bullet Performs color conversion (`COLOR_BGR2RGB`) prior to MediaPipe ingestion and renders live landmark overlays for caregiver visual feedback.

4.3 MediaPipe Face Landmarker Integration

- bullet Uses the `mediapipe.tasks.python.vision.FaceLandmarker` API with the `face_landmarker.task` bundle.
- bullet Configured in **IMAGE** mode with face presence/detection confidence thresholds set to 0.5 .
- bullet Returns 478 3D landmark points in normalized coordinate space ($x, y, z \in [0.0, 1.0]$). Fallback logic (`_compute_gaze_from_eyes`) uses eye corner midpoints if sub-pixel iris landmarks are unavailable.

4.4 Mathematical Computation: NumPy

- bullet Calculates Euclidean norm distance vectors (`np.linalg.norm`) for eye widths and heights.
- bullet Computes EAR vector products for blink detection.
- bullet Executes statistical IQR filtering (`np.mean`, P_{25} , P_{75}).
- bullet Applies Exponential Moving Average (EMA) smoothing to eliminate high-frequency eye tremor:

$$x_{\text{smooth}}(t) = \alpha \cdot x_{\text{raw}}(t) + (1 - \alpha) \cdot x_{\text{smooth}}(t-1), \quad \alpha = 0.25$$

4.5 Text-to-Speech: pytttsx3

- bullet Provides 100% offline, cross-platform speech synthesis.
- bullet Executed inside a dedicated background thread (**TTSWorker(QThread)**) to ensure the GUI main loop remains smooth during speech playback.
- bullet Configurable speech rate (80–250 WPM) and voice profile selection via the Settings panel.

4.6 Word Prediction Engine Architecture

A hybrid dual-layer prediction architecture:

1. **Local Frequency Dictionary (prediction/predictor.py)**: Loads **dictionary.json** containing frequency-sorted English vocabulary prioritized for common AAC needs (medical, comfort, personal, social). Performs instant prefix matching ($<1\text{ms}$ latency).
2. **Cloud LLM Prediction Engine**: Asynchronously queries the Groq API (**llama-3.3-70b-versatile**, temperature=0.3) in a background daemon thread. Constructs prompt context using full sentence history to yield context-aware next-word and completion candidates.

4.7 Context-Aware Abbreviation Expander

- bullet Monitored by **prediction/abbreviation_expander.py**.
- bullet Heuristics detect typed short tokens (lge 2 characters) with low vowel ratios (e.g. **"ftk"**, **"brb"**, **"pls"**).
- bullet Combines the caretaker's recent question/topic with prior conversation history into a structured prompt sent to Groq Llama-3.3 70B.
- bullet Parses JSON suggestions (e.g. [{"**expanded_text**": **"from the kitchen"**, **"confidence"**: **0.85**}]) and updates the suggestion strip without clearing user input.

4.8 Quick Phrases Engine

- bullet Pre-loaded dictionary (**PHRASE_CATEGORIES**) organized into 5 operational domains: *Urgent*, *Medical*, *Social*, *Quick Response*, and *Comfort*.
- bullet One-gaze selection immediately sends phrase text to **pytttsx3** for vocalization.

4.9 Caregiver Mode & Settings Infrastructure

- bullet **Caregiver Input**: Caretakers can type questions into **ContextBar** or use microphone input powered by AssemblyAI real-time streaming STT.
- bullet **Mouse & Keyboard Shortcuts**:
- bullet Full mouse-click support on all key widgets.
- bullet Keyboard shortcuts: **F5** (Recalibrate), **F11** (Toggle Fullscreen), **Esc** (Exit Fullscreen).
- bullet **Configuration Persistence**: Saves horizontal calibration parameters to **~/.gazespeak/calibration.json**.

5. Evaluation and Results

5.1 Evaluation Framework & Metrics

To evaluate GazeSpeak academically, a structured multi-metric framework is proposed:

1. Typing Speed & Efficiency:

- bullet **Words Per Minute (WPM)**:
$$\text{WPM} = \frac{\text{Characters}}{5} \times \frac{60}{\text{Time (sec)}}$$

bullet **Keystroke Savings (KSS)**: Percentage of typed characters saved by prediction/abbreviation:

$$\text{KSS} = \left(1 - \frac{\text{Actual Selection Steps}}{\text{Total Characters in Final Text}} \right) \times 100\%$$

2. Accuracy & Stability Metrics:

bullet **Fixation Accuracy / Target Selection Error Rate**: Ratio of unintended key selections to total selections.

bullet **Calibration Span & Repeatability**: Measurement of horizontal ratio variance across sessions.

3. Emergency Alert Performance:

bullet **Blink Alert Sensitivity & Specificity**: Precision of true positive SOS triggers vs false positives during normal blink activity.

4. User Satisfaction & Usability:

bullet **System Usability Scale (SUS)**: 10-item standard usability questionnaire.

bullet **NASA-TLX (Task Load Index)**: Measures mental demand, physical demand, and temporal demand.

5.2 System Performance Benchmarks

Based on software architecture measurements:

bullet **Iris Tracking Processing Latency**: ~15–25 ms per frame on standard CPU (30 fps webcam rate).

bullet **Local Prediction Latency**: $< 1 \text{ ms}$ (instantaneous).

bullet **Groq LLM Prediction Latency**: ~300–500 ms (asynchronous, non-blocking UI).

bullet **Dwell Time Range**: 300 ms – 2500 ms (default target: 800 ms).

bullet **Typing Rate Projection**:

bullet Raw eye-gaze typing: ~2.5–4.0 WPM.

bullet With Word Prediction & Abbreviation Expansion: ~6.0–12.0 WPM.

bullet With Binary Conversation Mode: ~15.0–25.0 WPM equivalent speech output.

5.3 System Limitations

1. **Lighting Dependencies**: Standard RGB webcams struggle in extremely dim environments ($< 20 \text{ lux}$) compared to dedicated infrared cameras.

2. **Head Pose Shift**: Significant head movement (outside a $\approx 15^\circ$ box) shifts eye landmarks relative to the webcam, requiring recalibration (**F5**).

3. **Iris Occlusion**: Severe ptosis (drooping eyelids) common in late-stage neuromuscular conditions can reduce landmark detection confidence.

4. **Network Dependencies**: Advanced features (Groq LLM, AssemblyAI STT, Twilio SMS) require an active internet connection, though local core typing and offline TTS remain fully operational offline.

6. Discussion and Future Work

6.1 Broader Implications for Assistive Technology

GazeSpeak demonstrates that high-performance, low-latency AAC solutions can be implemented using ubiquitous hardware (laptops/webcams) and modern vision models. By replacing complex 2D pointing with 1D discrete step navigation and leveraging generative LLMs for natural conversation narrowing, GazeSpeak bridges the financial and usability gap for ALS patients globally.

6.2 Future Research & System Expansion

- 1. Head Pose Compensation (Perspective-n-Point / PnP):** Integrating 3D head pose estimation to dynamically adjust gaze calibration when the patient moves their head.
- 2. Vertical Gaze Navigation:** Expanding the discrete horizontal step engine to include vertical gaze zones (LOOK UP / LOOK DOWN) for multi-directional menu control.
- 3. Edge / On-Device LLMs:** Replacing cloud Groq API calls with local quantized models (e.g. LLaMA 3 8B via Ollama / ONNX) to enable 100% offline context awareness.
- 4. Multi-Language AAC Support:** Extending dictionary datasets and LLM system prompts to support multi-lingual ALS patients.
- 5. Smart Home / IoT Integration:** Connecting selection actions to Home Assistant MQTT/REST endpoints to allow patients to control lighting, thermostats, and bed positioning.

6.3 Clinical Validation Roadmap

- bullet Conducting formal clinical trials with ALS patients in rehabilitation centers.
- bullet Long-term longitudinal studies evaluating eye fatigue (electrooculography / NASA-TLX) over extended daily usage compared to commercial Tobii hardware.